

# Stat 427/627 Statistical Machine Learning

## In-class Lab 8: Principal Components and Partial Least Squares

### Contents

1	Concepts of principal component	2
2	Principal Components regression (PCR)	8
3	Partial Least Squares regression (PLS)	10

Recall the Auto data set in the ISLR2 package. This data frame has 392 observations on the following 9 variables.

- mpg: miles per gallon
- cylinders: Number of cylinders between 4 and 8
- displacement: Engine displacement (cu. inches)
- horsepower: Engine horsepower
- weight: Vehicle weight (lbs.)
- acceleration: Time to accelerate from 0 to 60 mph (sec.)
- year: Model year (modulo 100)
- origin: Origin of car (1. American, 2. European, 3. Japanese)
- name: Vehicle name

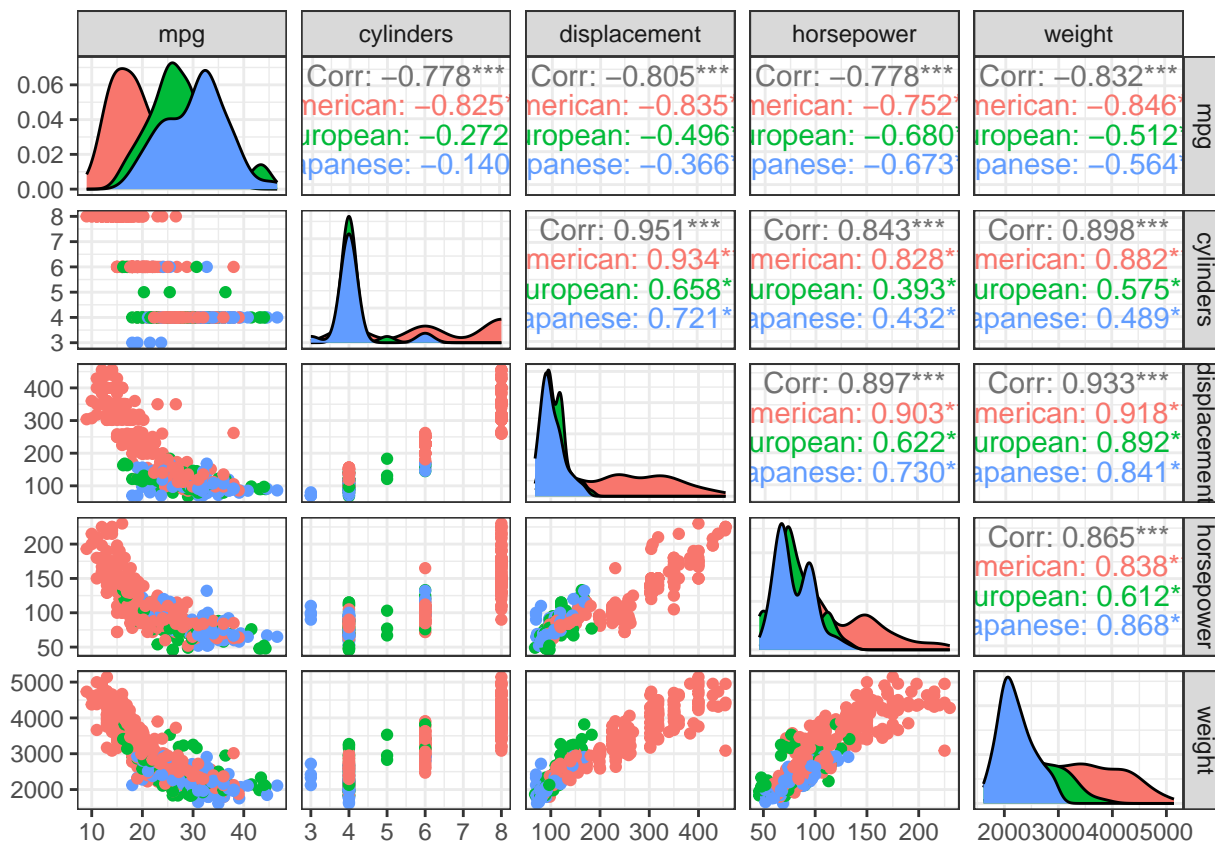
```
library(ISLR2)
colnames(Auto)

## [1] "mpg"          "cylinders"    "displacement" "horsepower"   "weight"
## [6] "acceleration" "year"        "origin"       "name"

auto.data <- Auto
auto.data$country <- factor(auto.data$origin, labels = c("American", "European",
                                                         "Japanese"))

library("ggplot2")
library("GGally")

ggpairs(auto.data, columns = 1:5, ggplot2::aes(colour=country)) + theme_bw()
```



## 1 Concepts of principal component

If we use a linear regression model to predict mpg, we'll have high correlation amongst the predictors.

```
mpg.lmF <- lm(mpg ~ . - origin - name, data=auto.data)
mpg.lmF
```

```
##
## Call:
## lm(formula = mpg ~ . - origin - name, data = auto.data)
##
## Coefficients:
## (Intercept)      cylinders      displacement      horsepower
##      -17.95460        -0.48971         0.02398        -0.01818
##      weight      acceleration          year countryEuropean
##      -0.00671         0.07910         0.77703         2.63000
## countryJapanese
##         2.85323
```

```
car::vif(mpg.lmF)
```

```
##           GVIF Df GVIF^(1/(2*Df))
## cylinders 10.73771 1      3.276854
## displacement 22.93795 1      4.789358
## horsepower  9.957265 1      3.155513
## weight     11.074349 1      3.327814
## acceleration 2.625906 1      1.620465
```

```
## year          1.301373  1          1.140777
## country       2.096060  2          1.203236
```

We will focus on the predictors now.

```
auto.X <- model.matrix(mpg.lmF)[, -1] # Design matrix
```

Function `prcomp()` conducts the Principal Component Analysis.

*# Caution: the following line of code is not ideal. It is to illustrate an issue.*

```
auto.pc <- prcomp(auto.X)
summary(auto.pc)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  855.6585 38.90971 16.16207 3.31353 1.697 0.5249 0.4167
## Proportion of Variance  0.9976  0.00206  0.00036 0.00001 0.000 0.0000 0.0000
## Cumulative Proportion  0.9976  0.99962  0.99998 1.00000 1.000 1.0000 1.0000
##              PC8
## Standard deviation   0.2446
## Proportion of Variance 0.0000
## Cumulative Proportion 1.0000
```

```
auto.pc
```

```
## Standard deviations (1, ..., p=8):
```

```
## [1] 855.6585163  38.9097121  16.1620689   3.3135262   1.6966834   0.5249057
## [7]  0.4167494   0.2446327
```

```
##
```

```
## Rotation (n x k) = (8 x 8):
```

```
##              PC1      PC2      PC3      PC4
## cylinders    -0.0017926225 -0.0133245279  0.007294275  0.001414710
## displacement -0.1143412856 -0.9457785881  0.303312504 -0.009143349
## horsepower   -0.0389670412 -0.2982553337 -0.948761071 -0.043076559
## weight       -0.9926735354  0.1207516411  0.002454212  0.001480458
## acceleration  0.0013528348  0.0348264293  0.077006895  0.059516278
## year         0.0013368415  0.0238516081  0.042819254 -0.996935229
## countryEuropean 0.0001308250  0.0024889942 -0.002857670  0.022100094
## countryJapanese 0.0002103564  0.0003765828 -0.004796684 -0.012089823
##              PC5      PC6      PC7      PC8
## cylinders    -1.719368e-02  0.9911554803  0.1211162208  4.909265e-02
## displacement  1.059355e-02 -0.0146594359 -0.0006512752 -4.394368e-03
## horsepower    8.646402e-02  0.0038232742  0.0034425206  4.435100e-03
## weight       -3.152970e-03 -0.0002093216 -0.0003053766 -5.729471e-06
## acceleration  9.944974e-01  0.0168319859  0.0012233398  1.799780e-03
## year         5.549653e-02 -0.0001647840  0.0240346554 -7.643176e-03
## countryEuropean 9.052576e-05 -0.0483462982  0.6888706846 -7.229226e-01
## countryJapanese 1.150938e-03  0.1214929883 -0.7142804151 -6.891098e-01
```

We see that the 1st principal component contains a huge portion of the total variation of X variables, and it is dominated by variable “weight”. Looking at the data, we see that weight simply has the largest values.

For this reason, the variables are usually standardized first (subtract each variable by its mean, then divide each centered-variable by its standard deviation). This is done using `scale=TRUE` argument.

```
auto.pcs <- prcomp(auto.X, scale=TRUE)
summary(auto.pcs)
```

```
## Importance of components:
```

```
##          PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  2.133 1.1356 0.9497 0.80460 0.62678 0.35401 0.24235
## Proportion of Variance 0.569 0.1612 0.1127 0.08092 0.04911 0.01567 0.00734
## Cumulative Proportion 0.569 0.7302 0.8429 0.92383 0.97294 0.98861 0.99595
##          PC8
## Standard deviation  0.18005
## Proportion of Variance 0.00405
## Cumulative Proportion 1.00000
```

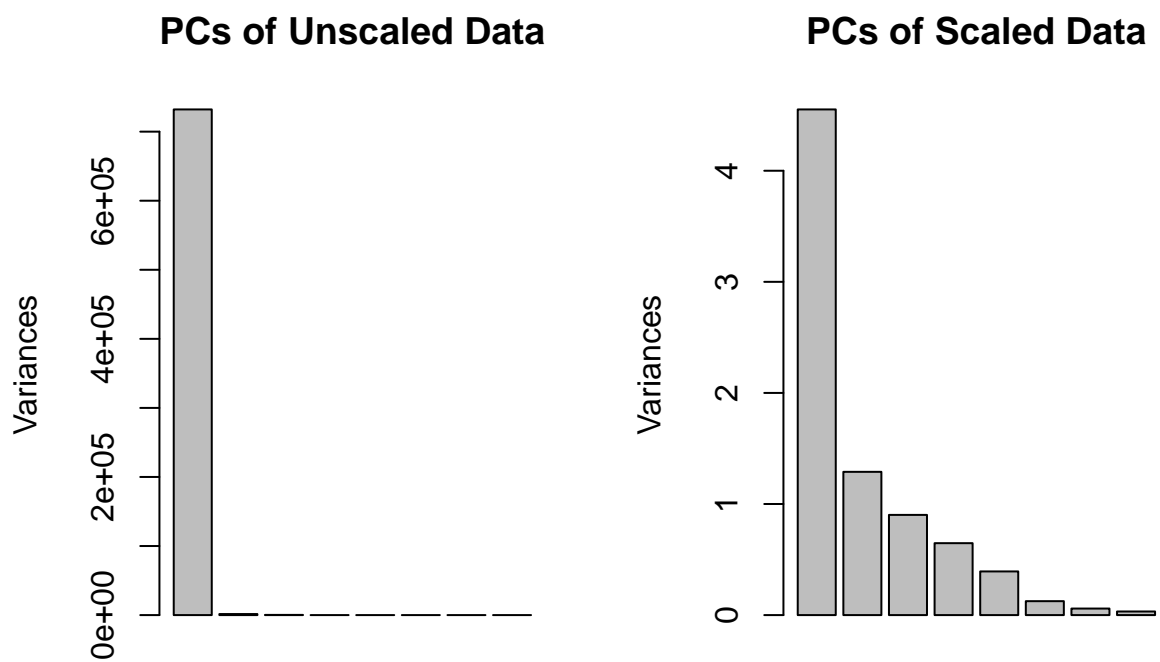
```
round(cor(auto.pcs$x), 4)
```

```
##      PC1 PC2 PC3 PC4 PC5 PC6 PC7 PC8
## PC1   1  0  0  0  0  0  0  0
## PC2   0  1  0  0  0  0  0  0
## PC3   0  0  1  0  0  0  0  0
## PC4   0  0  0  1  0  0  0  0
## PC5   0  0  0  0  1  0  0  0
## PC6   0  0  0  0  0  1  0  0
## PC7   0  0  0  0  0  0  1  0
## PC8   0  0  0  0  0  0  0  1
```

```
auto.pcs$rotation # Extract the loading (i.e., transformation)
```

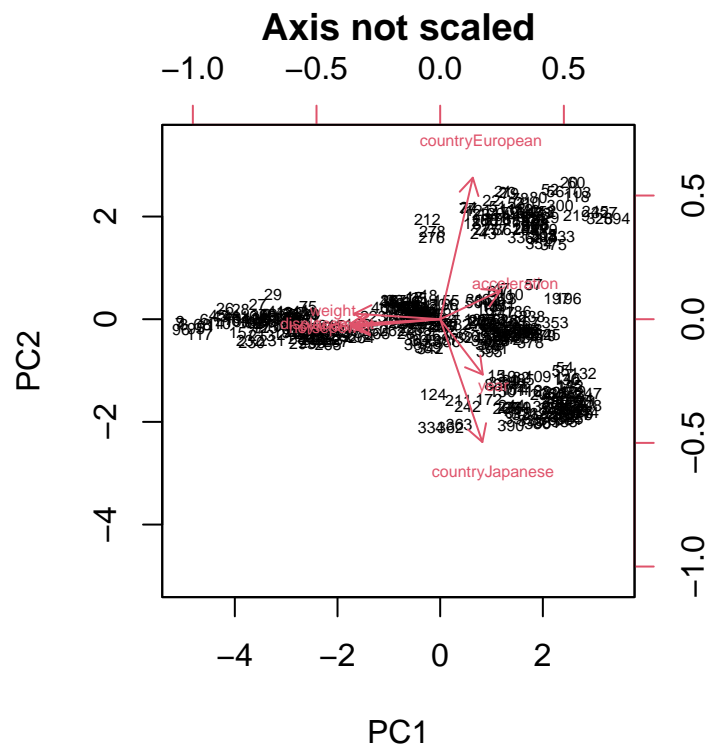
```
##          PC1          PC2          PC3          PC4          PC5
## cylinders    -0.4418880 -0.03183661  0.1513759  0.08112451 -0.17557152
## displacement -0.4568839 -0.02468263  0.1334604  0.05932358 -0.09729606
## horsepower   -0.4427613 -0.04007041 -0.1065626 -0.15129036 -0.20758380
## weight       -0.4343646  0.02796626  0.2382679  0.10749689 -0.31123269
## acceleration  0.3038014  0.14422734  0.4574875  0.70453506 -0.30582654
## year         0.2153593 -0.27979875  0.7122813 -0.58611507 -0.13985800
## countryEuropean 0.1648074  0.71513286 -0.1544922 -0.33648172 -0.55320193
## countryJapanese 0.2131795 -0.62087485 -0.3881227  0.04831876 -0.63295244
##          PC6          PC7          PC8
## cylinders    -0.714889530  0.219309935  0.42922755
## displacement -0.157462297  0.046158763 -0.85611061
## horsepower    0.571133678  0.610447413  0.16180753
## weight        0.318542871 -0.705651981  0.21433158
## acceleration  0.127257355  0.266640825 -0.02042327
## year         0.007589381  0.067111659 -0.01832329
## countryEuropean -0.114101557 -0.006045755 -0.07244053
## countryJapanese -0.084877251 -0.059530689 -0.06868871
```

```
par(mfrow=c(1, 2))
plot(auto.pc, main="PCs of Unscaled Data")
screplot(auto.pcs, main="PCs of Scaled Data")
```




A biplot shows the first 2 PCs scores (bottom, left) and their loadings (top, right).

```
biplot(auto.pcs, scale=0, cex=0.5, main="Axis not scaled") # Loadings are plotted as is.
```



```
biplot(auto.pcs, scale=1, cex=0.5, mai="Scaled axis") # Loadings are scaled.
```



For observations in the data frame, their PC scores are saved. PC scores for new observations can be computed using `predict()` function.

```
auto.pcs$x[c(1:3), ]
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## 1 -2.631256  0.17038496 -1.033093   0.327091974  0.4948202 -0.7335556 -0.07630174
## 2 -3.373310  0.10389029 -1.105045   0.110134107  0.2522577 -0.2311124  0.29240679
## 3 -2.984687  0.09245312 -1.259333 -0.009259703  0.5125006 -0.5249749  0.20558186
##          PC8
## 1  0.07875985
## 2 -0.07450685
## 3  0.06308631
```

```
newx <- data.frame(cylinders = 6, displacement = 250, horsepower = 88,
                   weight = 3021, acceleration = 17, year = 73,
                   countryEuropean = 0, countryJapanese = 0)
predict(auto.pcs, newdata = newx)
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,] -0.4081881  0.2825017  0.1067126  1.103147  0.4889963 -0.3771085 -0.08551976
##          PC8
## [1,] -0.3085454
```

## 2 Principal Components regression (PCR)

It will be more convenient to use `pcr()` function from package `pls` to use PCs in regression analysis. It puts 2 analysis into 1 function:

- First, conduct PCA on the predictors and save the PC scores for all observations.
- Next, conduct linear regression use the PCs as the predictors in the model.

```
# install.packages("pls")
library(pls)
```

### 2.1 Fit PCR

Be sure to include `scale=TRUE` argument.

```
auto.pcr <- pcr(mpg ~ .-name=origin, data=auto.data, scale=T)
summary(auto.pcr)
```

```
## Data:      X dimension: 392 8
## Y dimension: 392 1
## Fit method: svdpc
## Number of components considered: 8
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          56.9   73.02   84.29   92.38   97.29   98.86   99.59  100.00
## mpg        71.8   73.64   73.96   79.25   79.25   80.22   81.55   82.42
```

- What is the value of  $R^2$  if we use:
  - The first PC for the regression?
  - The first 2 PCs for the regression?
  - The first 5 PCs for the regression?
  - All PCs (8 in this case) for the regression?

### 2.2 Cross-validation

`pcr()` function has built-in option for K-fold (default  $K = 10$ ) and LOOCV. Use argument `validation = "CV"` or `validation = "LOO"`.


```
auto.pcrCV <- pcr(mpg ~ .-name=origin, data=auto.data, scale=T, validation="CV")
summary(auto.pcrCV)
```

```
## Data:      X dimension: 392 8
## Y dimension: 392 1
## Fit method: svdpc
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              7.815   4.164   4.034   4.036   3.597   3.612   3.538
## adjCV           7.815   4.163   4.032   4.035   3.594   3.608   3.535
##      7 comps  8 comps
## CV          3.430   3.358
## adjCV        3.426   3.353
##
```



```
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X      56.9   73.02   84.29   92.38   97.29   98.86   99.59  100.00
## mpg     71.8   73.64   73.96   79.25   79.25   80.22   81.55   82.42
```


```
validationplot(auto.pcrCV)
```



```
auto.pcrL00 <- pcr(mpg ~ .-name-origin, data=auto.data, scale=T, validation="L00")
summary(auto.pcrL00)
```

```
## Data:      X dimension: 392 8
## Y dimension: 392 1
## Fit method: svdpc
## Number of components considered: 8
##
## VALIDATION: RMSEP
## Cross-validated using 392 leave-one-out segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV           7.815  4.163  4.038  4.031  3.61  3.621  3.545
## adjCV        7.815  4.163  4.038  4.031  3.61  3.621  3.545
##      7 comps 8 comps
## CV       3.431  3.361
## adjCV    3.431  3.360
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X      56.9   73.02   84.29   92.38   97.29   98.86   99.59  100.00
```

```
## mpg      71.8    73.64    73.96    79.25    79.25    80.22    81.55    82.42
validationplot(auto.pcrL00)
```



### 2.3 Predict the response for new data.

`predict()` function works for `pcr()` output object. Use argument `ncomp` = to declare how many PC's to use in the model.

```
predict(auto.pcr, newdata=auto.data[c(1, 2), ], ncomp=4) # If 4 PCs are used.
```

```
## , , 4 comps
##
##      mpg
## 1 13.92093
## 2 12.13321
```

## 3 Partial Least Squares regression (PLS)


Use function `plsrv()`. It is similar to `pcr()`

```
auto.plsCV <- plsrv(mpg ~ .-name-origin, data=auto.data, scale=T, validation="CV")
summary(auto.plsCV)
```

```
## Data:      X dimension: 392 8
## Y dimension: 392 1
## Fit method: kernelpls
## Number of components considered: 8
```

```
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept) 1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## CV          7.815   3.993   3.600   3.531   3.422   3.391   3.357
## adjCV       7.815   3.991   3.597   3.527   3.416   3.387   3.352
##      7 comps 8 comps
## CV      3.360   3.360
## adjCV    3.356   3.355
##
## TRAINING: % variance explained
##      1 comps 2 comps 3 comps 4 comps 5 comps 6 comps 7 comps 8 comps
## X      56.73   68.84   80.75   84.08   93.48   94.88   99.33  100.00
## mpg     74.32   79.37   80.29   81.71   82.00   82.35   82.38   82.42
```

```
validationplot(auto.plsCV)
```



```
names(auto.plsCV)
```

```
## [1] "coefficients" "scores" "loadings" "loading.weights"
## [5] "Yscores" "Yloadings" "projection" "Xmeans"
## [9] "Ymeans" "fitted.values" "residuals" "Xvar"
## [13] "Xtotvar" "fit.time" "ncomp" "method"
## [17] "center" "scale" "validation" "call"
## [21] "terms" "model"
```

```
predict(auto.plsCV, newdata=auto.data[c(1, 2), ], ncomp=4) # If 4 PCs are used.
```

```
## , , 4 comps
##
##      mpg
## 1 15.92016
## 2 13.97132
```