

# hw\_04

lisa liubovich

2024-06-11

## 1. College Applications

Predict the number of applications received based on the other variables in the College data set in {ISLR2}.

Use the following code to split the data set randomly. Use the subset college.data for the most of the analysis. Set aside holdout.data data until the last 2 sub-questions.

```
library(ISLR2)
data("College")
my.college <- College[-484, ] # remove an extreme case.
#my.college <- College[College$Apps <=16000, ] # remove several extreme case.
train.pct <- 0.78
set.seed(2024)
Z <- sample(nrow(my.college), floor(train.pct*nrow(my.college)))
college.data <- my.college[Z, ]
holdout.data <- my.college[-Z, ]
```

Recall that we fit the data in homework 1.

- we fit the **full** model with all 17 predictors. That is:

```
college.lmF <- lm(Apps ~ ., data = college.data)
college.lmF
```

```
##
## Call:
## lm(formula = Apps ~ ., data = college.data)
##
## Coefficients:
## (Intercept)    PrivateYes        Accept        Enroll    Top10perc    Top25perc
## -4.685e+02    -5.008e+02    1.397e+00   -5.142e-01    5.189e+01   -1.480e+01
## F.Undergrad  P.Undergrad    Outstate    Room.Board        Books        Personal
##  6.690e-02    3.188e-02   -6.002e-02    2.188e-01    1.534e-01   -2.228e-03
##          PhD          Terminal    S.F.Ratio  perc.alumni        Expend        Grad.Rate
## -7.742e+00   -4.247e+00    1.057e+01   -5.725e+00    6.035e-02    6.194e+00
```

Using stepwise selection, AIC is the smallest with 12 predictors:

Apps ~ Private + Accept + Enroll + Top10perc + Top25perc + F.Undergrad + Outstate + Room.Board + PhD + perc.alumni + Expend + Grad.Rate

Using the best-subset algorithm. BIC is the smallest with 7 predictors:

Apps ~ Private + Accept + Top10perc + Outstate + Room.Board + PhD + Expend

Continue working on this data set and the above models.

- (a) Compute the variance inflation factor and comment on the severity of the collinearity of the data. Why is “collinearity” a concern, even if the model is correct?

- (a) Some of the variables like Accept, Enroll, and F.Undergrad have VIF values over 5 or 10, indicating that they suffer from more severe multicollinearity. Even if this model is correct, multicollinearity is a concern because it increases the variance and thus the standard errors, which make for less accurate coefficient estimates, inferences, and overly wide confidence intervals.

```
library(car)
```

```
## Loading required package: carData
```

```
vif_college <- vif(college.lmF)
print(vif_college)
```

```
##      Private      Accept      Enroll      Top10perc      Top25perc      F.Undergrad
##      2.865245     10.150354     21.012832      6.881086      5.392519     16.472473
## P.Undergrad      Outstate      Room.Board      Books      Personal      PhD
##      1.698963      4.379916      2.101790      1.142078      1.363831      4.002334
##      Terminal      S.F.Ratio      perc.alumni      Expend      Grad.Rate
##      3.917552      1.982753      1.950469      2.964559      1.929511
```

- (b) Evaluate prediction accuracy of your selected models based on AIC and BIC. Estimate the prediction mean squared error by 10-fold cross-validation. Recall that `glm( ..., family=gaussian)` fits linear regression and its outcome can be used in `cv.glm()` (boot package) for cross-validation.

```
library(boot)
```

```
##
```

```
## Attaching package: 'boot'
```

```
## The following object is masked from 'package:car':
```

```
##
```

```
##      logit
```

```
AIC_college <- glm(Apps ~ Private + Accept + Enroll + Top10perc + Top25perc + F.Undergrad + Outstate + Room.Board + PhD + Expend, data=college.data)
```

```
BIC_college <- glm(Apps ~ Private + Accept + Top10perc + Outstate + Room.Board + PhD + Expend, data=college.data)
```

```
cv.glm(college.data, glmfit = AIC_college, K = 10)$delta[1]
```

```
## [1] 997934.5
```

```
cv.glm(college.data, BIC_college, K = 10)$delta[1]
```

```
## [1] 980591.9
```

The BIC model has a lower predicted MSE by 10 fold CV.

- (c) Consider the **full** model with all 17 predictors (reminder: use data frame college.data you recreated at the beginning). Use functions in package glmnet to fit a ridge regression.

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
college.X <- model.matrix(college.lmF)[, -1]  
dim(college.X)
```

```
## [1] 605 17
```

- Select  $\lambda$  chosen by (default 10-fold) cross-validation.

```
college.ridgeCV <- cv.glmnet(x=college.X, y = college.data$Apps, alpha=0,  
                             lambda=seq(10, 0, by=-0.1))  
college.ridgeCV
```

```
##
```

```
## Call: cv.glmnet(x = college.X, y = college.data$Apps, lambda = seq(10, 0, by = -0.1), alpha = 0)
```

```
##
```


```
## Measure: Mean-Squared Error
```

```
##
```

```
##      Lambda Index Measure      SE Nonzero  
## min      8.6      15 1005968 204425      17  
## 1se     10.0       1 1006195 204191      17
```

- Plot the results of the cross-validation.

```
plot(college.ridgeCV)
```



- Report the estimated MSE of the model based on your selected  $\lambda$

```
college.ridgeCV
```

```
##
## Call: cv.glmnet(x = college.X, y = college.data$Apps, lambda = seq(10,      0, by = -0.1), alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      8.6      15 1005968 204425      17
## 1se     10.0       1 1006195 204191      17
```

estimated MSE is 1005968.

- (d) Consider the **full** model with all 17 predictors (reminder: use data frame college.data you recreated at the beginning). Use functions in package glmnet to fit a LASSO regression.


```
college.LASSOCV <- cv.glmnet(x=college.X, y = college.data$Apps, alpha=1,
                             lambda=seq(10, 0, by=-0.1))
college.LASSOCV
```

```
##
## Call: cv.glmnet(x = college.X, y = college.data$Apps, lambda = seq(10,      0, by = -0.1), alpha = 1)
##
```

```
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      10      1  992888 139144      16
## 1se      10      1  992888 139144      16
```

- Select  $\lambda$  chosen by (default 10-fold) cross-validation.
- Plot the results of the cross-validation.

```
plot(college.LASSOCV)
```



- Report the estimated MSE of the model based on your selected  $\lambda$ .

```
college.LASSOCV
```

```
##
## Call:  cv.glmnet(x = college.X, y = college.data$Apps, lambda = seq(10,      0, by = -0.1), alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      10      1  992888 139144      16
## 1se      10      1  992888 139144      16
```

estimated MSE is 992888

- (e) Fit a PCR model on college.data, with  $M$  (the number of principal components) chosen by cross validation. Prepare a validation plot. Report the estimated test error (MSE), along with the value of  $M$  selected by cross-validation.

```
library(pls)


##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##      loadings

pcr_collegeCV <- pcr(Apps ~ ., data = college.data, scale = "TRUE", validation = "CV")
summary(pcr_collegeCV)

## Data:      X dimension: 605 17
## Y dimension: 605 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           3459    3456    1508    1461    1440    1177    1096
## adjCV        3459    3456    1507    1460    1438    1175    1095
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1097    1100    1100    1068    1047    1045    1042
## adjCV        1095    1097    1098    1065    1044    1043    1040
##      14 comps 15 comps 16 comps 17 comps
## CV           1030    1011    1013    1005
## adjCV        1028    1008    1010    1002
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X       50.8029  87.69   95.72   97.84   98.79   99.47   99.92   99.97
## Apps    0.6885  81.35   82.47   83.16   89.10   90.46   90.52   90.59
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X       100.00  100.00  100.00  100.00  100.00  100.00  100.00
## Apps    90.61   91.18   91.55   91.59   91.66   91.88   92.34
##      16 comps 17 comps
## X       100.00  100.00
## Apps    92.36   92.51

validationplot(pcr_collegeCV)
```



estimated MSE =  $\text{RMSEP}^2 = 1013^2 = 1026169$


- (f) Fit a PLS (partial least squares) model on college.data, with  $M$  (the number of principal components) chosen by cross validation. Prepare a validation plot. Report the estimated test error (MSE), along with the value of  $M$  selected by cross-validation.

```
college_plsCV <- plsr(Apps ~ ., data = college.data, scale = "TRUE", validation = "CV")
summary(college_plsCV)
```

```
## Data:      X dimension: 605 17
## Y dimension: 605 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           3459    1466    1425    1323    1116    1095    1095
## adjCV        3459    1465    1429    1324    1115    1093    1093
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1101    1103    1082    1047    1026    1013    1009
## adjCV        1099    1100    1085    1045    1024    1010    1006
##      14 comps 15 comps 16 comps 17 comps
## CV           1007    1007    1008    1000.1
## adjCV        1005    1005    1005     997.1
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
```

```
## X      37.09   79.99   94.31   96.92   98.72   99.44   99.75   99.97
## Apps   82.12   83.20   85.88   90.03   90.50   90.55   90.59   90.64
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X      100.00  100.00  100.00  100.00  100.00  100.00  100.00  100.00
## Apps   90.75   91.62   91.97   92.31   92.35   92.36   92.36
##      16 comps 17 comps
## X      100.00  100.00
## Apps   92.37   92.51
```

```
validationplot(college_plsCV)
```



estimated MSE =  $1008^2 = 1016064$

- (g) Summarize and comment on the results obtained from the following models. Recommend a model, and justify your choice.

Method	Number of predictors	Estimated Prediction MSE
Least Squares 1: model with the smallest AIC	12	997934.5
Least Squares 2: model with the smallest BIC	7	980591.9
Ridge Regression (lambda.min)	17	1005968
Lasso (lambda.min)	16	992888
Lasso (lambda.1se)	16	992888
PCR PLS	15	1016064

Based on the summarized results, the model with the smallest BIC (Least Squares 2), which includes 7 predictors, is recommended. This model achieves an estimated prediction MSE of 980,591.9, indicating



superior predictive performance compared to the other models evaluated. BIC's preference for parsimony aligns well here, emphasizing a balance between model complexity and predictive accuracy. While ridge regression and lasso offer regularization and feature selection benefits, the chosen least squares model provides a straightforward and effective choice, balancing predictive power with model simplicity. Adjustments may be considered based on specific criteria such as interpretability or computational efficiency, but for overall predictive performance based on the given metrics, Least Squares 2 stands out as the optimal choice.

- (h) Apply the above models to the hold-out data `holdout.data` that we created at the beginning. Which model wins this contest in terms of prediction accuracy? (This should be the first time you use observations in `holdout.data` data frame.)

```
X_holdout <- as.matrix(holdout.data[, -1])
y_holdout <- holdout.data$Apps

AIC_college_2 <- glm(Apps ~ Private + Accept + Enroll + Top10perc + Top25perc + F.Undergrad + Outstate +
BIC_college_2 <- glm(Apps ~ Private + Accept + Top10perc + Outstate + Room.Board + PhD + Expend, data =
cv.glm(holdout.data, glmfit = AIC_college_2, K = 10)$delta[1]

## [1] 1335918

cv.glm(holdout.data, BIC_college_2, K = 10)$delta[1]

## [1] 1420287

college.ridgeCV2 <- cv.glmnet(x = X_holdout, y = y_holdout, alpha = 0, lambda = seq(10, 0, by = -0.1))
college.ridgeCV2

##
## Call: cv.glmnet(x = X_holdout, y = y_holdout, lambda = seq(10, 0, by = -0.1),      alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      0.0   101   74.35   67.64      17
## 1se      0.4    97  138.80  121.63      17

college.LASSOCV2 <- cv.glmnet(x = X_holdout, y = y_holdout, alpha = 1, lambda = seq(10, 0, by = -0.1))
college.LASSOCV2

##
## Call: cv.glmnet(x = X_holdout, y = y_holdout, lambda = seq(10, 0, by = -0.1),      alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min         0   101 3.627e-25 9.777e-26      17
## 1se         0   101 3.627e-25 9.777e-26      17
```

```
college_plsCV2 <- plsr(Apps ~ ., data = holdout.data, scale = "TRUE", validation = "CV")
summary(college_plsCV2)
```

```
## Data:      X dimension: 171 17
## Y dimension: 171 1
## Fit method: kernelppls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3739    1433    1389    1338    1214    1226    1235
## adjCV           3739    1423    1275    1321    1205    1215    1224
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1257    1267    1262    1232    1233    1233    1231
## adjCV        1244    1254    1253    1218    1218    1218    1217
##      14 comps 15 comps 16 comps 17 comps
## CV          1258    1269    1276    1250
## adjCV        1241    1252    1258    1232
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          45.42   50.88   89.12   97.29   98.32   99.31   99.47   99.96
## Apps       86.94   90.53   91.00   91.99   92.19   92.20   92.27   92.28
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          100.00  100.00  100.00  100.00  100.00  100.00  100.00
## Apps       92.35   93.12   93.27   93.36   93.43   93.48   93.52
##      16 comps 17 comps
## X          100.00  100.00
## Apps       93.53   93.9
```

- (h) AIC MSE = 1398325, BIC MSE = 1365042, ridge MSE = 21.69, LASSO MSE =  $3.817 \cdot 10^{-25}$ , PLS MSE =  $1228^2 = 1507984$ . LASSO wins the contest by far.
- (i) **Stat-627** Compare your estimated prediction MSE from the training data college.data (part i) and the resulting MSE from the holdout.data (part j). Is there anything “surprising” that worth investigation? If yes, what are the possible causes? (Note. It is not surprising to see a tuned “best” model not to perform the best on the testing data.)

Yes, the prediction MSEs for ridge regression and LASSO are considerably smaller than all the other prediction MSEs as well as all of the training MSEs. This is for a couple reasons: first, the shrinkage methods provide a better bias-variance tradeoff by shrinking the coefficients (thus reducing bias) and significantly reducing variance. This is consistent with the fact that shrinkage methods in general tend to make models less complex by shrinking the coefficients to almost 0 (ridge) or even 0 (LASSO). The smaller prediction MSEs observed indicate that these regularization techniques have effectively optimized the model complexity for improved prediction accuracy on the holdout data. The MSEs are also so much smaller because of the use of cross validation, where the lambda parameter (regularization strength) is selected to minimize prediction error on validation data. This ensures that the model is tuned to perform well on unseen data, further contributing to the smaller prediction MSEs observed.